Tutorial II.III - DataFrames in Julia

Optimization with Julia



Introduction

Imagine a DataFrame as a digital spreadsheet. It's a way to organize and work with data in rows and columns. Each column can hold different types of information, like names, ages, or salaries. In this tutorial, we'll learn how to create DataFrames, add and change data, and perform simple operations on our data.

> **Note**
> Before we start, make sure you have the DataFrames package installed. If you're not sure how to do this, check the previous tutorial on package management!

Let's begin by importing the DataFrames package:

```
# Import the DataFrames package
using DataFrames
```

Section 1 - Creating DataFrames

A DataFrame in Julia is akin to a table in SQL or a spreadsheet - each column can have its own type, making it highly versatile. A DataFrame can be created using the DataFrame constructor and passing key-value pairs where the key is the column name and the value is an array of data. For more help, use `?` in the REPL and type `DataFrame`. Example:

```
students = DataFrame(
    Name = ["Elio", "Bob", "Yola"],
    Age = [18, 25, 29],
)
```

```
3×2 DataFrame
 Row │ Name    Age
     │ String  Int64
─────┼────────────────
   1 │ Elio       18
   2 │ Bob        25
   3 │ Yola       29
```

Exercise 1.1 - Create a DataFrame

Create and Test a DataFrame. Create a DataFrame named `employees` with the columns `Name`, `Age`, and `Salary`, and populate it with the specified data: John is `28` years old and earns `50000`. Mike is `23` years old and earns `62000`. Frank is `37` years old and earns `90000`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert employees == DataFrame(
    Name = ["John", "Mike", "Frank"],
    Age = [28, 23, 37],
    Salary = [50000, 62000, 90000]
) "Check the column names, their order, and the values - they should be Name,
Age and Salary with the data from the task."
println("DataFrame created successfully!")
println(employees)
```

> **Tip**
> Remember, for more help, use `?` in the REPL and type `DataFrame`.

Section 2 - Accessing and Modifying Data

Accessing columns in a DataFrame can be done using the dot syntax, while rows can be accessed via indexing. Modification of data is straightforward; just assign a new value to the desired cell. To access the column `Name` in our DataFrame `employees`, we could do:

```
employees.Name
```

```
3-element Vector{String}:
 "John"
 "Mike"
 "Frank"
```

To access the third name specifically, we could do:

```
employees.Name[3]
```

```
"Frank"
```

Exercise 2.1 - Access the `Age` Column

Access the `Age` column from the DataFrame and save it in a new variable `ages`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert ages == [28, 23, 37]
println("Correct, the Ages column is: ", ages)
```

Exercise 2.2 - Update John's Salary

Update John's salary to `59000`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert employees.Salary[1] == 59000
println("Modified DataFrame: ")
println(employees)
```

Section 3 - Filtering Data

Logical indexing can be used to filter rows in a DataFrame based on conditions. To filter the DataFrame to include only employees named "Frank" we could do:

```
all_franks = employees[employees.Name .== "Frank", :]
```

```
1×3 DataFrame
 Row │ Name    Age    Salary
     │ String  Int64  Int64
─────┼──────────────────────
   1 │ Frank      37   90000
```

Alternatively, the filter function provides a powerful tool to extract subsets of data based on a condition:

```
all_franks = filter(row -> row.Name == "Frank", employees)
```

```
1×3 DataFrame
 Row │ Name    Age    Salary
     │ String  Int64  Int64
─────┼──────────────────────
   1 │ Frank      37   90000
```

Exercise 3.1 - Filter the DataFrame

Filter the DataFrame to include only employees with salaries above `60000`. Save the resulting employees in the DataFrame `high_earners`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert nrow(high_earners) == 2
println("High earners: ")
println(high_earners)
```

Section 4 - Basic Data Manipulation

Julia provides functions for basic data manipulation tasks, such as sorting DataFrames. The `sort` function can be used to order the rows in a DataFrame based on the values in one or more columns. To see how to use the function, type `?` into the REPL (terminal) and type `sort`. DataFrames can also be grouped (with `groupby` and `combine`) and joined - you will meet these functions later in the course when we prepare data for optimization models.

Exercise 4.1 - Sort the DataFrame

Sort the DataFrame based on the `Age` column and save it as `sorted_df`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert sorted_df.Age[1] == 23
println("DataFrame sorted by age: ")
println(sorted_df)
```

Section 5 - Loop over DataFrames

Sometimes, you might need to iterate over the rows of a DataFrame to perform operations on each row individually. Julia provides a convenient way to do this using the `eachrow` function. For example, if we want to check for each employee if they have a salary above 60000, we can do the following:

```
for row in eachrow(employees)
    if row.Salary > 60000
        println("$(row.Name) earns more than 60000")
```

```
    end
end
```

```
Mike earns more than 60000
Frank earns more than 60000
```

Here, `row` is a `DataFrameRow` - a view into the DataFrame. This means that if you change a value through `row`, the change is written back to the `employees` DataFrame. We can access the values of a column by using the dot syntax. To create a new column, assign a vector (or broadcast a single value with `.=`) to a new column name. For example, to create a new column called `VacationDays` in the `employees` DataFrame, we can do one of the following:

```
employees.VacationDays = zeros(Int, nrow(employees))
employees.VacationDays .= 0
```

```
3-element Vector{Int64}:
 0
 0
 0
```

Here, `nrow` returns the number of rows in a DataFrame and `zeros(Int, n)` creates a vector of `n` zeros.

Exercise 5.1 - Loop over DataFrame

Create a new column called `Bonus` in the `employees` DataFrame. The bonus should be calculated as 10% of the salary for employees over 30, and 5% for those 30 and under. Use a loop to iterate over the rows and calculate the bonus.

> ! **Important**
> This exercise assumes that you updated John's salary to `59000` in Exercise 2.2. If you skipped that exercise, go back and complete it first - otherwise the test below will fail.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert filter(
    row -> row.Bonus == 2950,
    employees
).Name == ["John"] "John should have a bonus of 2950"
@assert filter(
    row -> row.Bonus == 3100,
    employees
).Name == ["Mike"] "Mike should have a bonus of 3100"
@assert filter(
    row -> row.Bonus == 9000,
```

```
    employees
).Name == ["Frank"] "Frank should have a bonus of 9000"
println("Great job! All the bonuses are correct!")
```

Section 6 - Filling a new DataFrame with values

Do you remember the `push!` function? We can use it to fill a new DataFrame with values row by row. This pattern is worth remembering: later in the course, you will use it to collect the results of your optimization models in a DataFrame. For example, we can create a new DataFrame called `WorkingHours` and fill it with the values from the `employees` DataFrame. Imagine that the company has a policy, where employees above 30 work 30 hours a week, and employees under 30 work 40 hours a week:

```
# Create a new DataFrame
WorkingHours = DataFrame(
    Name = String[],
    Hours = Int[]
)
# Loop over the rows of the employees DataFrame
for row in eachrow(employees)
    if row.Age < 30
        push!(WorkingHours, (
            Name = row.Name,
            Hours = 40
        ))
    else
        push!(WorkingHours, (
            Name = row.Name,
            Hours = 30
        ))
    end
end
println(WorkingHours)
```

```
3×2 DataFrame
 Row │ Name    Hours
     │ String  Int64
─────┼────────────────
   1 │ John       40
   2 │ Mike       40
   3 │ Frank      30
```

Exercise 6.1 - Fill a new DataFrame with `push!`

Now practice this pattern yourself. Create a new DataFrame called `Overtime` with the columns `Name` (a `String` column) and `ExtraHours` (an `Int` column). Then, loop over the rows of the `employees` DataFrame with `eachrow` and `push!` one row per employee: employees with a salary above `60000` get `5` extra hours, all others get `10`.

```
# YOUR CODE BELOW
```

```
# Test your answer
@assert nrow(Overtime) == 3 "Overtime should have one row per employee, thus 3
rows."
@assert Overtime.Name == ["John", "Mike", "Frank"] "The names should appear in
the same order as in employees."
@assert Overtime.ExtraHours == [10, 5, 5] "John should have 10 extra hours,
Mike and Frank 5 each."
println("Perfect, you filled a new DataFrame row by row!")
```

Conclusion

Fantastic work! You've completed the tutorial on DataFrames in Julia. You've seen how to create DataFrames and access, modify and filter data. Continue to the next file to learn more.

Solutions

You will find the solutions to all exercises online here in the `solutions` folders for each part. However, I strongly encourage you to work on these exercises independently without searching explicitly for the exact answers to the exercises. Understanding someone else's solution is very different from developing your own. Use the lecture notes and try to solve the exercises on your own. This approach will significantly enhance your learning and problem-solving skills.

Remember, the goal is not just to complete the exercises, but to understand the concepts and improve your programming abilities. If you encounter difficulties, review the lecture materials, experiment with different approaches, and don't hesitate to ask for clarification during class discussions.